

Business Proposal

Forget-lah

SP-ARK Agents | Team code: N63VHYEX

**Intelligent patient follow-up.
Staff remain in control.**

A practical agentic workflow for the work that begins after a reminder is sent.

The problem: follow-up takes more than a message

Dental clinics maintain follow-up schedules for routine examinations, preventive care and ongoing treatment. Staff repeatedly review records, identify who needs attention, contact patients and resolve the replies. As records grow, consistent follow-up becomes harder, creating the risk of missed appointments and delayed treatment.

Staff must interpret doctor instructions, check prerequisites and available options, update the correct system, and escalate exceptions. This repetitive coordination competes with critical clinic work. Patients can forget appointments or tests, misunderstand preparation instructions, face language barriers or need clarification after clinic hours. Older patients may need additional support.

The proposed solution

Forget-lah connects existing clinic information, patient communication and staff oversight into one evidence-backed follow-up workflow. It uses AI where understanding is useful, deterministic controls where certainty is required, and clinic staff where human judgment belongs.


Dental follow-up is the origin, not the limit. The dental, scan-prerequisite and ophthalmology demonstrations are representative examples of a common workflow across clinical settings.

Submission scope

An AWS-hosted hackathon prototype with synthetic clinic records, controlled WhatsApp test exchanges and finished scenario demonstrations. The product is an intelligent follow-up layer: it is not a chatbot, an EMR, a general clinic scheduler or a medical-diagnosis system.

Repository: github.com/satworkz/forget-lah

How Forget-lah fits the clinic



Two routes, one follow-up experience

For API-backed clinics, an API is the secure software connection used to read approved information and write permitted updates. The existing Clinic System remains the master record. Forget-lah does not replace it.

For clinics without a modern API, Intelligent Intake understands existing spreadsheet exports. Uncertain rows go to staff for correction; approval is required before import. Forget-lah then manages only the imported follow-up lifecycle in its own PostgreSQL data store. Availability must come from the clinic source or staff-supplied options.

Patients and staff

WhatsApp is the implemented hackathon channel. Enrolled test patients can reply or initiate follow-up questions, including outside clinic hours when the service is running. Answers must be grounded in available clinic instructions; clinical or unresolved questions go to staff. Sandbox delivery rules still apply. SMS, a patient app and voice interaction are future channels.

The staff dashboard shows cases, replies, instructions, evidence and exceptions. Staff can review intake, make validated changes and accept responsibility for escalations. A queued handoff is not treated as owned until a named staff member accepts it.

From repetitive work to verified outcomes

The value is less repetitive coordination, so staff can focus on work that requires judgment.

Today: staff coordinate every step	With Forget-lah
Find patient and read notes	Detect eligible follow-up and load approved context.
Send reminder and interpret each reply	Coordinate through WhatsApp; interpret intent and clinic instructions.
Check prerequisite, availability and update	Validate policy and source evidence before permitted actions.
Confirm the outcome and record history	Acknowledge verified results; retain workflow evidence and audit records.
Handle exceptions among routine messages	Route clinical concern, uncertainty and unsupported requests to staff.

What the demonstration proves

Scenario	Capability demonstrated
Ahmad: dental follow-up	Attendance intent is separated from reported swelling/pain; clinical concern is handed to staff.
Priya: scan prerequisite	Doctor-defined instructions constrain rescheduling and eligible options; a patient selection is validated.
Alex: multilingual context	Chinese-language coordination reasons over the conflict between driving and the clinic instruction about blurred vision/accompaniment.
Intelligent Intake	Semantic mapping, REVIEW/READY, staff correction, approval and provenance.
Staff change + Security & Audit	Human intervention remains validated and traceable, with source and notification evidence.

Expected business value: fewer manual touches, more consistent follow-up, visibility of exceptions and a practical route for both API and legacy-data clinics. These benefits require pilot measurement; the synthetic demo is not evidence of clinical outcomes, production savings or reduced no-show rates.

How we would measure a clinic pilot

Proposed stretch targets - not achieved production results.

Establish a one-week manual baseline, then review a four-week controlled pilot with clinic-approved case types. Compare similar cases and report sample size, exclusions and unresolved cases. Agree targets and escalation cover with the clinic before launch.

Measure	Proposed target	Measurement
Manual staff touches	At least 60% fewer per eligible routine case	Staff interventions per case, using workflow events and a short activity log. Review weekly.
Active coordination time	At least 50% reduction	Compare staff handling time for similar cases; exclude patient-response and appointment waiting time.
Routine completion without staff action	At least 60% of eligible routine cases	Verified administrative outcomes without staff action / all eligible routine cases. Clinical decisions are excluded.
Intake usefulness and quality	At least 70% READY without correction; at least 95% field accuracy	Compare with staff-labelled reference records. Report REVIEW rate and corrections; staff approval remains mandatory.
Exception routing and ownership	100% of predefined escalation tests routed correctly	Audit test outcomes, reason, named acceptance and time to acceptance. No handoff closes as owned without acceptance.

Measurement discipline

Track total time to confirmed, rescheduled or escalated state as a separate observation. A fast message is not a verified outcome. Record failed deliveries and source conflicts instead of dropping them from the analysis. Review the staff queue daily; review the measures weekly and at pilot close.

Economic assessment

Estimate staff hours saved per 100 comparable follow-ups from observed handling-time differences. Compare the value of those hours with hosting, messaging, model usage and operational support costs. No monetary saving or return on investment is claimed before measurement.

Delivery, controls and growth

Stage	Deliverable	Estimate / assumption
Week 1	Dental workflow, synthetic Clinic API, three agent roles, durable state and first staff escalation.	5 working days; synthetic records and agreed follow-up rules available.
Week 2	Multilingual and prerequisite workflows, controlled WhatsApp, validated changes and Intelligent Intake.	5 working days; gateway/test-channel access and scoped options available.
Final stage	Security/audit checks, failure scenarios, AWS demo, video and pilot measurement plan.	5 working days; prototype acceptance, not production certification.

The staged plan is a proposal estimate, not a retrospective record of effort. The repository and demo already provide the prototype evidence described in this submission.

Safe operation and manual fallback

Clinic staff approve imports, own escalations and explicitly initiate staff changes. Patient-requested changes require explicit selection or confirmation where applicable. Every action still passes authority, evidence and source checks. Clinical decisions remain outside automation.

Source or model failure leaves a visible unresolved case, not an invented success. Staff inspect current records and contact patients through clinic procedures. If notification fails after a verified update, retain the update and retry communication rather than repeating the appointment change.

Feasibility and responsible scale

Durable jobs, leases, source receipts and separate worker loops support continuity and bounded concurrency. Selective model calls reduce unnecessary inference. The system is built with high availability in mind, but the current single-host demo does not prove host failover or production throughput. Reactive wake-ups, redundancy, managed recovery and capacity testing are next deployment stages.

Future - outside the delivered prototype

Patient Companion with offline reminders/calendar support; voice accessibility; caregiver-authorised support; trusted identity and consent integration; Follow-up Pattern Intelligence; diagnostic/lab/physiotherapy integrations; and consent-based group reminders. Education reminders are exploratory reuse, not part of the clinical scope.

Basis and references

Prepared against the supplied SMYA Business Proposal Guidelines and the current repository: README; FEATURE_GUIDE; BRIDGE_INTAKE; SECURITY_ARCHITECTURE; PERFORMANCE; ORGANISER_GATEWAY; FAILURE_HANDLING; AWS_DEMO. Implementation detail and verification limits appear in the companion Technical Document. No production healthcare certification or real-patient outcome is claimed.